

Mastering Java performance with async-profiler

Andrei Pangin Odnoklassniki

2019

About me





- Principal Engineer @ Odnoklassniki
- JVM enthusiast
- Top #jvm answerer on Stack Overflow
- Author of async-profiler
 https://github.com/jvm-profiling-tools/async-profiler

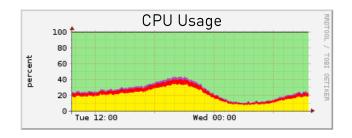
Agenda

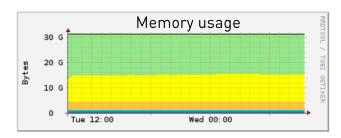


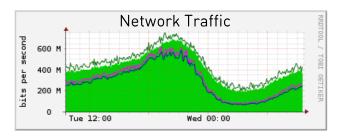
- 1. Profiling theory
- 2. What's wrong with the classical approach
- 3. Introduction to async-profiler
- 4. Allocation profiling
- 5. Advanced techniques

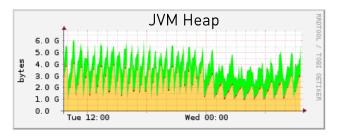
What to optimize?

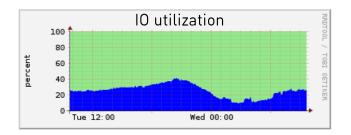


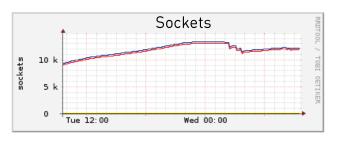






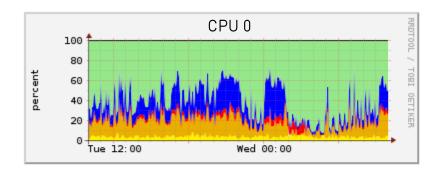


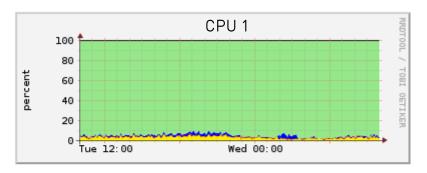


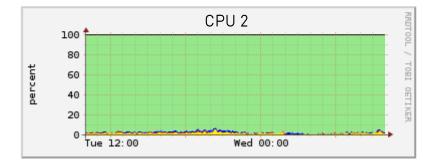


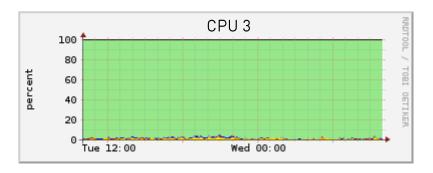
To profile or not to profile?













CPU profiling

How to profile



Instrumenting

Trace method transitions

Sloooow

Sampling

Periodic snapshots

Production ready

Thread Dump



Java API

Thread.getAllStackTraces()



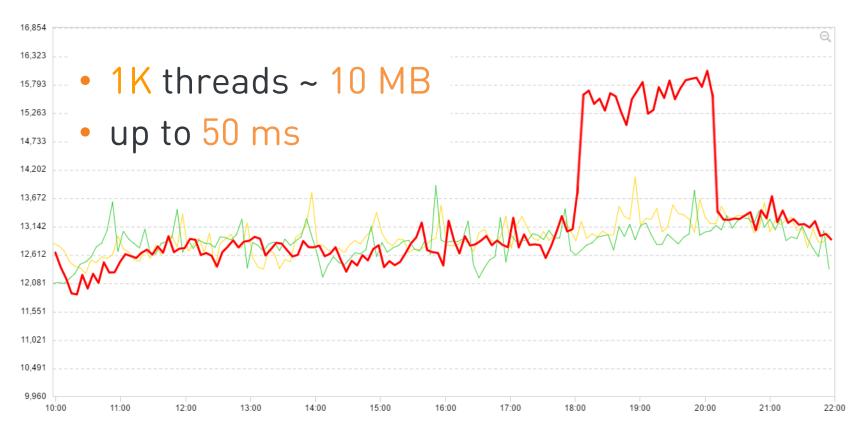
Map<Thread, StackTraceElement[]>

Native (JVMTI)

GetAllStackTraces()

Overhead





Advantages

8

- Simple
- All Java platforms
- No extra JVM options required

Supported by many profilers





Replying to @AndreiPangin

I'm happy with **Profiler ... no need for another tool

9:37 PM - 18 May 2018



(\)



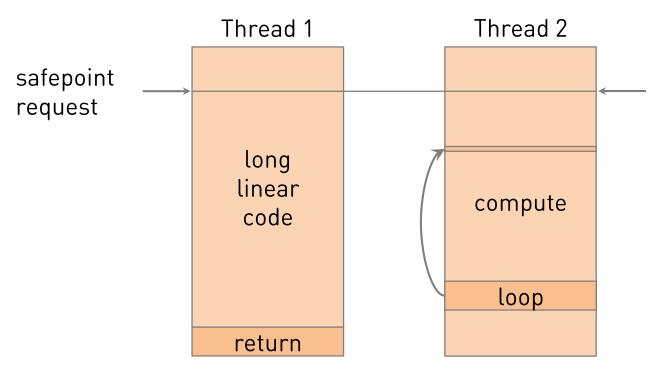




Demo

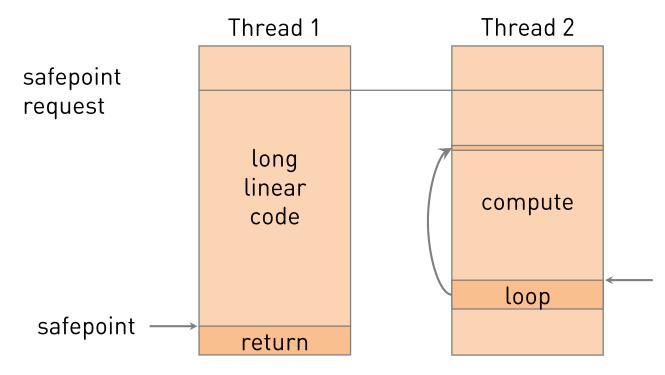
Safepoint





Safepoint







Why (Most) Sampling Java Profilers Are F Terrible

This post builds on the basis of a previous post on safepoints. If you've not read it you might feel lost and confused. If you have read it, and still feel lost and confused, and you are certain this feeling is related to the matter at hand (as opposed to an existential crisis), please ask away.

So, now that we've established what safepoints are, and that:

- 1. Safepoint polls are dispersed at fairly arbitrary points (depending on execution mode, mostly at uncounted loop back edge or method return/entry).
- 2. Bringing the JVM to a global safepoint is high cost





Why (Most) Sampling Java Profilers Are F Terrible

because of Safepoint bias ...not the profiler's fault...

Native methods

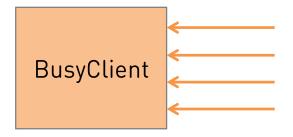


```
Socket s = new Socket(host, port);
InputStream in = s.getInputStream();
while (in.read(buf) >= 0) {
    // keep reading
}
```

Native methods



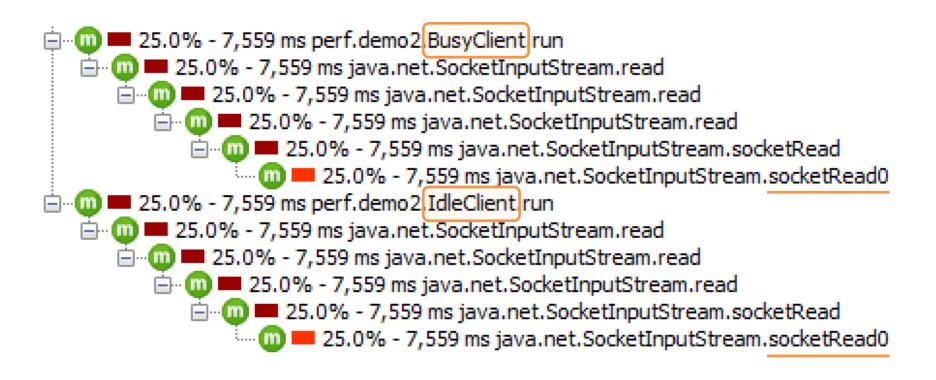
```
Socket s = new Socket(host, port);
InputStream in = s.getInputStream();
while (in.read(buf) >= 0) {
    // keep reading
}
```





Native methods





Solving GetAllStackTraces problems



- Avoid safepoint bias
- Skip idle threads
- Profile native code



AsyncGetCallTrace

How it works



```
AsyncGetCallTrace(ASGCT_CallTrace *trace,
                  jint depth,
                  void* ucontext)
```



from itimer signal handler

- Oracle Developer Studio
- github.com/jvm-profiling-tools/honest-profiler
- github.com/jvm-profiling-tools/async-profiler

honest-profiler example



Method	Total %	Self#% ▲
▼ Thread-1	0,20 %	0,00 %
perf.demo2 <mark>.ldleClient.run</mark>	0,20 %	0,20 %
▼ Thread-2	96,22 %	0,00 %
▼ perf.demo2.BusyClient.run	95,82 %	0,20 %
▼ java.net.SocketInputStream.read	95,62 %	0,00 %
▼ java.net.SocketInputStream.read	95,62 %	0,00 %
▼ java.net.SocketInputStream.read	95,62 %	0,60 %
▼ java.net.SocketInputStream.socketRead	94,62 %	0,00 %
java.net.SocketInputStream.socketRead0	94,62 %	94,62 %

AsyncGetCallTrace

8

- Pros
 - Only active threads
 - No safepoint bias
 - -XX:+DebugNonSafepoints
- Cons
 - * No Windows support
 - * No native or JVM code

The bad part



Method	Total % ▼	Self#%
▼ main	98,07 %	0,00 %
AGCT.UnknownJavaErr5	95,22 %	95,22 %
▼ perf.demo1.ArrayListTest.main	2,64 %	0,29 %
sun.launcher.LauncherHelper.checkAndLoadMain	0,21 %	0,00 %
▼ Unknown Thread(s)	1,85 %	0,00 %
AGCT.UnknownNotJava3	1,85 %	1,85 %

The bad part



```
double avg(Number... numbers) {
    double sum = 0;
    for (Number n : numbers) {
        sum += n.doubleValue();
    return sum / numbers.length;
x = avg(123, 45.67, 890L, 33.3f, 999, 787878L);
```

The bad part



Method	Total % ▼	Self#%
▼ main	99,79 %	0,00 %
AGCT.UnknownJavaErr5	55,16 %	55,16 %
▼ perf.demo3.Numbers.main	44, <mark>42 %</mark>	0,00 %
▼ perf.demo3.Numbers.loop	44, <mark>42 %</mark>	10,67 %
▼ perf.demo3.Numbers.avg	33,75 %	24,35 %
java.lang.Long.doubleValue	6,88 %	6,88 %
java.lang.Integer.doubleValue	1,33 %	1,33 %
java.lang.Float.doubleValue	0,70 %	0,70 %
java.lang.Double.doubleValue	0,49 %	0,49 %

Problems



```
enum {
  ticks no Java frame
                               = 0,
  ticks_no_class_load
                               = -1,
  ticks GC active
                               = -2,
  ticks unknown not Java
                               = -3,
  ticks not walkable not Java = -4,
  ticks unknown Java
                               = -5,
  ticks_not_walkable_Java
                               = -6,
  ticks unknown state
                               = -7,
  ticks thread exit
                               = -8,
  ticks deopt
                               = -9,
  ticks_safepoint
                               = -10
                                           src/share/vm/prims/forte.cpp
```

AsyncGetCallTrace fails to traverse valid Java stacks

Status:



Details

4 P4 R

OPEN

tbd

Priority:

Affects Version/s:

8u121, 9, 10

Resolution: Unresolved

Component/s: hotspot

Labels: AsyncGetCallTrace

analyzer

Fix Version/s:

Subcomponent: svc

CPU: x86 64

People

Assignee:

Reporter:

Andrei Pangin

Unassigned

Votes:

• Vote for this issue

Watchers:

3 Start watching this issue

Description

There is a number of cases (observed in real applications) when profiling with AsyncGetCallTrace is useless due to HotSpot inability to walk through Java stack. Here is the analysis of such cases.

#1. An application performs many System.arraycopy() and spends a lot of time inside

Dates

Created:

2017-04-06 17:32



Perf Events

PMU





Hardware counters

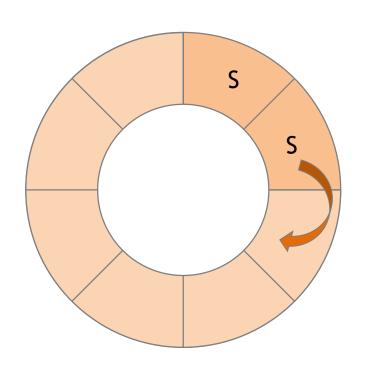
- Cycles, instructions
- Cache misses
- Branch misses
- etc.

perf_event_open

9

- Linux syscall
 - Subscribe to HW/OS events

- Samples
 - pid, tid
 - CPU registers
 - Call chain (user + kernel)



perf



```
$ perf record -F 1009 java ...
$ perf report
```

perf.wiki.kernel.org/index.php/Tutorial

perf



```
$ perf record -F 1009 java ...
$ perf report
  4.70%
           java
                    [kernel.kallsyms]
                                           [k] clear page c
                    libpthread-2.17.so
                                           [.] pthread_cond_wait
  2.10%
           java
  1.97%
           java
                    libjvm.so
                                           [.] Unsafe Park
                                           [.] Parker::park
  1.40%
           java
                    libjvm.so
  1.31%
           java
                    [kernel.kallsyms]
                                           [k] try to wake up
                                               0x00007f8510e9e757
  1.31%
           java
                    perf-18762.map
                    perf-18762.map
  1.21%
           java
                                               0x00007f8510e9e89e
  1.17%
           java
                    perf-18762.map
                                           [.] 0x00007f8510e9cc17
```

perf.wiki.kernel.org/index.php/Tutorial

perf-map-agent



github.com/jvm-profiling-tools/perf-map-agent

```
$ java -agentpath:/usr/lib/libperfmap.so ...
```

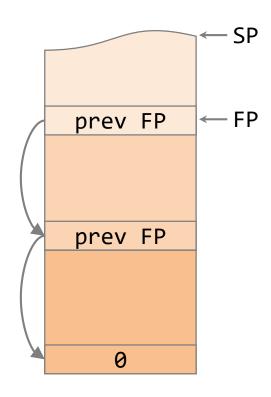


```
7fe0e91175e0 140 java.lang.String::hashCode
7fe0e9117900 20 java.lang.Math::min
7fe0e9117ae0 60 java.lang.String::length
7fe0e9117d20 180 java.lang.String::indexOf
```

JVM support for perf



-XX:+PreserveFramePointer



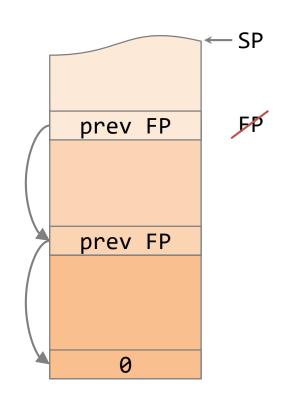
JVM support for perf



-XX:+PreserveFramePointer

• JDK ≥ 8u60

Permanent overhead (< 5%)



perf + FlameGraph



- 1. perf record
- 2. perf script
- 3. FlameGraph/stackcollapse-perf.pl
- 4. FlameGraph/flamegraph.pl



github.com/brendangregg/FlameGraph

Flame Graph



ip_local_out
ip_queue_xmit
tcp_transmit_skb

tcp_write_xmit

_tcp_push_pending_fra..

java_start start_thread iava-339

GC.. JavaThread::run

Flat profile



Hot Spots – Method	Self Time [%]
sun.nio.ch.FileDispatcherImpl.read0[native] ()	
java.net.SocketInputStream. socketRead0[native] ()	
one.nio.net.NativeSelector.epollWait[native] ()	
one.nio.net.NativeSocket.accept0[native] ()	
$sun.nio.ch. Server Socket Channel Impl. {\color{red} accept 0 [native]}\ ()$	
sun.management.ThreadImpl.dumpThreads0[native] ()	
sun.misc.Unsafe.unpark[native] ()	
sun.reflect.Reflection. getCallerClass[native] ()	
sun.nio.ch.NativeThread.current[native] ()	
sun.nio.ch.FileDispatcherImpl.write0[native] ()	
sun.misc.Unsafe.park[native] ()	
$org. a pache. cass and ra. locator. Dynamic Endpoint Snitch. {\bf get Datacenter}\ ()$	
$org. a pache. cass and ra. locator. Abstract Network Topology Snitch. {\color{red} compare Endpoints}\ ()$	
java.lang.Object. hashCode[native] ()	
java.lang.Object.notifyAll[native] ()	
java.util.concurrent.ConcurrentHashMap. get ()	
java.security.AccessController.doPrivileged[native] ()	
java.lang.String. intern[native] ()	

Tree view

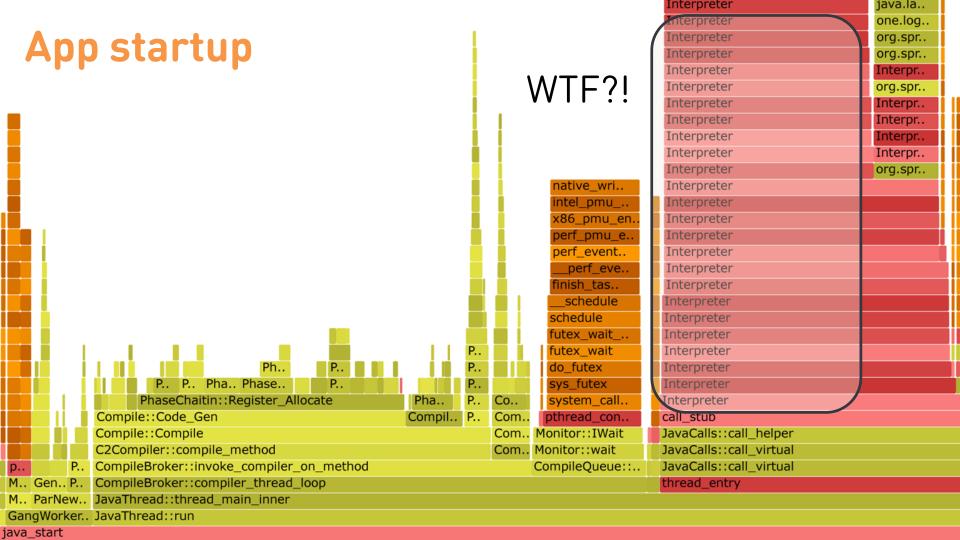


Call Tree - Method

- ⁷ Thread-6
 - arg.apache.cassandra.net.IncomingTcpConnection.run ()
 - org.apache.cassandra.net.IncomingTcpConnection.handleModernVersion ()
 - java.io.DataInputStream.readInt ()
 - \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 \(\)
 - java.io.BufferedInputStream.read ()
 - java.io.BufferedInputStream.fill ()
 - Sun.nio.ch.ChannelInputStream.read ()
 - Sun.nio.ch.SocketAdaptor\$SocketInputStream.read ()
 - sun.nio.ch.SocketChannelImpl.read ()
 - sun.nio.ch.IOUtil.read ()
 - sun.nio.ch.IOUtil.readIntoNativeBuffer ()
 - sun.nio.ch.SocketDispatcher.read ()
 - Sun.nio.ch.FileDispatcherImpl.read0[native] ()



Demo



Perf disadvantages



- No interpreted frames
- perf-map-agent for symbols
- -XX:+PreserveFramePointer
- JDK ≥ 8u60
- /proc/sys/kernel/perf_event_paranoid
- /proc/sys/kernel/perf_event_max_stack = 127
- «Big Data»

Mixed approach





perf_event_open

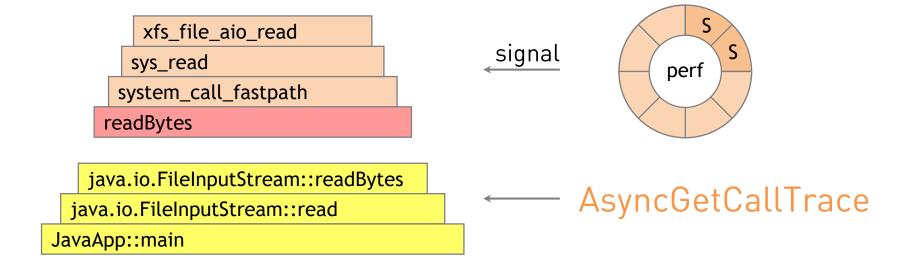
- Kernel + native stacks
- HW counters

- Entire Java
- Fast and precise

AsyncGetCallTrace

How it works in async-profiler



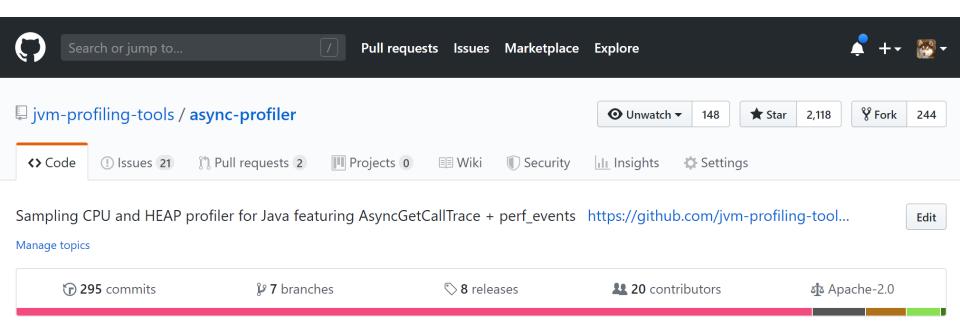




Introducing async-profiler

github.com/jvm-profiling-tools/async-profiler





github.com/jvm-profiling-tools/async-profiler



Download

Latest release (1.6):

- Linux x64 (glibc): async-profiler-1.6-linux-x64.tar.gz
- Linux x64 (musl): async-profiler-1.6-linux-x64-musl.tar.gz
- Linux ARM: async-profiler-1.6-linux-arm.tar.gz
- macOS x64: async-profiler-1.6-macos-x64.tar.gz

Previous releases

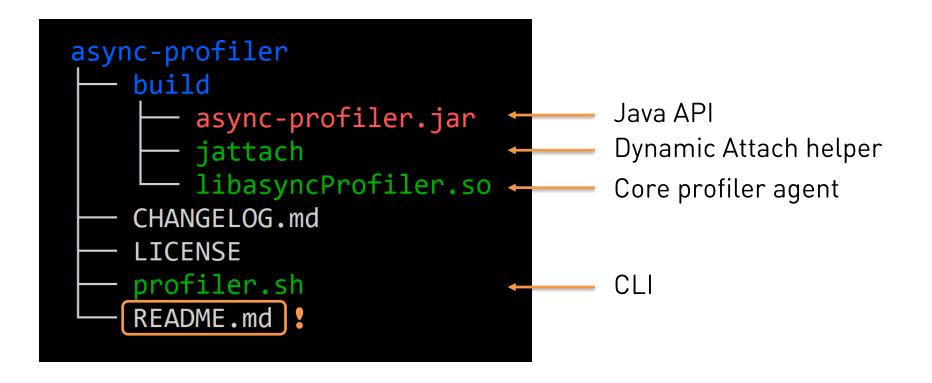
Supported platforms

- Linux / x64 / x86 / ARM / AArch64
- macOS / x64

Note: macOS profiling is limited to user space code only.

Package contents







Demo

Case study: file reading



```
byte[] buf = new byte[bufSize];
try (FileInputStream in = new FileInputStream(fileName)) {
    int bytesRead;
    while ((bytesRead = in.read(buf)) > 0) {
                                                Buffer size?
                                                   □ 64 K
                                                   □ 1 M
                                                   4 M
                                                   □ 16 M
                                                   □ 32 M
```

Down to Linux Kernel



```
clear_pag.. get..
                                                                      alloc_pages_no..
                                                                                                   lru..
                                                                                         me..
                                                                  alloc_pages_vma
                                                                                         me..
                                                                                                page..
                                                              handle mm fault
                                                              do_page_fault
                                                           do page fault
   .. copy_user_generic_string
                                    file rea..
                                             page fault
 generic file aio read
 xfs_file_aio_read
 do sync read
 vfs read
 sys read
 system call fastpath
(/usr/lib64/libpthread-2.17.so)
                                                                                                             _memmove_ssse3_
readBytes
java.io.FileInputStream::readBytes
java.io.FileInputStream::read
ReadFile::readFile
```

Down to Linux Kernel



Buffer size: 32M => 31M

```
__.. copy_user_generic_string
generic_file_aio_read

xfs_file_aio_read

do_sync_read

vfs_read

sys_read

system_call_fastpath

(/usr/lib64/libpthread-2.17.so)
    __memmove_ssse3_back

readBytes

java.io.FileInputStream::readBytes
java.io.FileInputStream::read

ReadFile::readFile
```

Case study: nanoTime



System.nanoTime() latency: $40 \text{ ns} \rightarrow 10 000 \text{ ns}$



"good" nanoTime

```
read_hpet
ktime_get_ts64
__.. posix_ktime_get_ts
sys_clock_gettime
s.. system_call_fast_compare_end
__vdso_clock_gettime
__clock_gettime
os::javaTimeNanos()
perf/Nanotime.measureTimeObjects
perf/Nanotime.main
```

"bad" nanoTime

Case study: nanoTime



System.nanoTime() latency: $40 \text{ ns} \rightarrow 10 000 \text{ ns}$

```
read_hpet
                                                                                    posix_ktime_get_ts
                                                                             sys_clock_gettime
                            [vdso]
                                                                          s.. system_call_fast_compare_end
                vdso_clock_gettime
                                                                         vdso clock gettime
          clock_gettime
                                                                        clock_gettime
    os::javaTimeNanos()
                                                                     os::javaTimeNanos()
perf/Nanotime.measureTimeObjects
                                                                    perf/Nanotime.measureTimeObjects
perf/Nanotime.main
                                                                    perf/Nanotime.main
               "good" nanoTime
                                                                                    "bad" nanoTime
```

```
[ 0.065868] TSC synchronization [CPU#0 -> CPU#1]:
[ 0.065870] Measured 679995254538 cycles TSC warp between CPUs, turning off TSC clock.
[ 0.065874] tsc: Marking TSC unstable due to check_tsc_sync_source failed
```





Replying to @AndreiPangin

I'm happy with **Profiler ... no need for another tool

9:37 PM - 18 May 2018



 \bigcirc







Yes

Yes

6+

No

Linux, macOS*

| | JFR | perf | async-profiler |
|------------|-----|----------------|----------------|
| Java stack | Yes | No interpreted | Yes |
| | | | |

Yes

Yes

8u60+

2-5%

Linux only

Possible

No

No

All

No

0

8*, 11+

Native stack

Kernel stack

JDK support

OS support

Permanent overhead

System-wide profiling



Wall clock profiling

Waste time doing nothing

8

- Thread.sleep
- Object.wait / Condition.await
- Wait to acquire a lock / semaphore / etc.
- Wait for I/O
 - Socket read
 - DB query
 - Disk I/O

Waste time doing nothing

8

- Thread.sleep
- Object.wait / Condition.await
- Wait to acquire a lock / semaphore / etc.
- Wait for I/O
 - Socket read
 - DB query
 - Disk I/O



Wall clock profiling



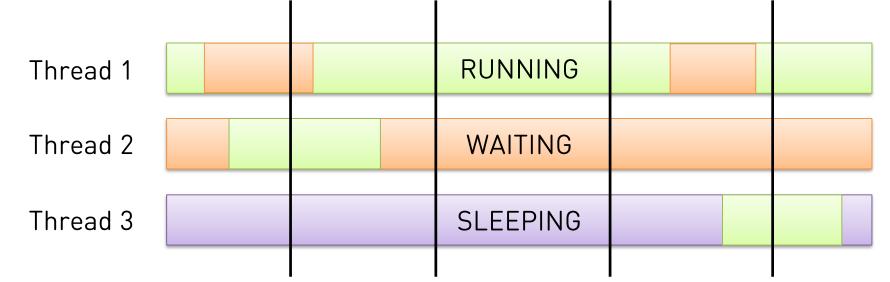




Wall clock profiling







Wall clock mode in async-profiler



| CLI option | Agent option |
|------------|--------------|
| -e wall | event=wall |
| -t | threads |
| -i 1ms | interval=1ms |

Lock contention





Profiler option

- -e lock
- -o svg=total

Lock contention





Profiler option

-e lock

-o svg=total



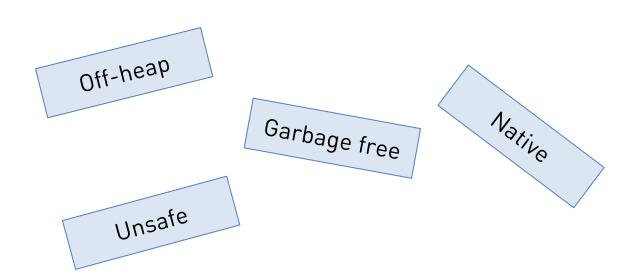
Demo



Allocation profiling

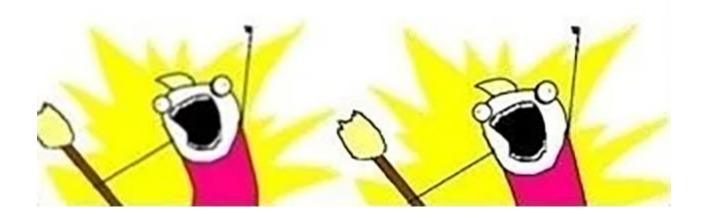


«new is the root of all evil»





We want full-featured Java!



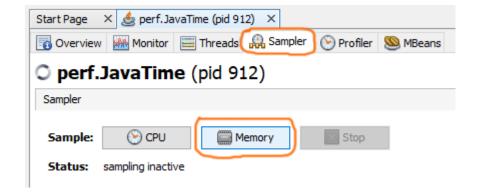
How many objects created?

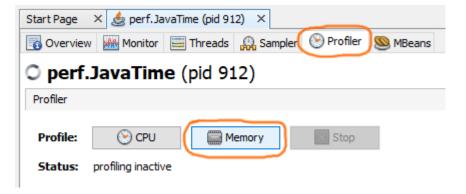


```
LocalDateTime deadline =
        LocalDateTime.now().plusSeconds(10);
while (LocalDateTime.now().isBefore(deadline)) {
   iterations++;
}
```

Profiling with VisualVM







Sampler



- Walk through heap
- Collect class histogram

| Class Name | Bytes [%] ▼ | Bytes | Instances | | |
|--|-------------|--------------------|-----------------|---|--|
| java.util.concurrent.ConcurrentHashMap | | 15,441,856 (20.7%) | 241,279 (12.2%) | • | |
| sun.util.calendar.ZoneInfo | | 13,510,112 (18.1%) | 241,252 (12.2%) | | |
| java.time.zone.ZoneRules | | 9,649,880 (12.9%) | 241,247 (12.2%) | | |
| java.time.LocalTime | | 5,790,672 (7.7%) | 241,278 (12.2%) | | |
| java.time.LocalDateTime | | 5,790,120 (7.7%) | 241,255 (12.2%) | | |
| java.time. LocalDate | | 5,790,120 (7.7%) | 241,255 (12.2%) | | |
| java.time.ZoneOffset[] | | 5,789,928 (7.7%) | 241,247 (12.2%) | | |
| java.time.ZoneRegion | | 5,789,928 (7.7%) | 241,247 (12.2%) | | |

Profiler



Bytecode instrumentation

```
→ Tracer.recordAllocation(ArrayList.class, 24);
List<String> list = new ArrayList<>();
```

getStackTrace() each Nth sample

DTrace / SystemTap

8

- -XX:+DTraceAllocProbes
- Slow path allocations

```
$ stap -e '
    probe hotspot.object_alloc { log(probestr) }
'
```

https://docs.oracle.com/javase/8/docs/technotes/guides/vm/dtrace.html https://epickrram.blogspot.ru/2017/09/heap-allocation-flamegraphs.html

Other tools



- Aprof
 https://code.devexperts.com/display/AProf
- Allocation Instrumenter
 https://github.com/google/allocation-instrumenter

Overhead



| Profiler | | x10 ⁶ ops | Slowdown |
|------------------------|----------------|----------------------|----------------------|
| No profiling | | 33,2 | |
| VisualVM Sampler | 1G
2G
4G | 28,2
25,6
18,1 | -15%
-23%
-45% |
| VisualVM Profiler | | 6,8 | -80% |
| JProfiler | | 2,8 | -92% |
| DTraceAllocProbes | | 13,3 | -60% |
| Aprof | | 18,9 | -43% |
| Allocation Instrumente | er | 22,8 | -31% |

JMC / Flight Recorder



| llocation Profile | | | |
|--|---------------------|-----------------|----------|
| Stack Trace | Average Object Size | Total TLAB size | Pressure |
| ✓ № perf.JavaTime.main(String[]) | 37 bytes | 22,75 GB | 99,94% |
| perf.JavaTime.measureTimeObjects(int) | 37 bytes | 22,75 GB | 99,94% |
| java.time.LocalDateTime.now() | 37 bytes | 22,75 GB | 99,94% |
| > java.time.Clock.systemDefaultZone() | 39 bytes | 19,33 GB | 84,919 |
| java.time.LocalDateTime.now(Clock) | 24 bytes | 3,42 GB | 15,029 |
| java.time.LocalDateTime.ofEpochSecond(long, int, ZoneOffset) | 24 bytes | 3,42 GB | 15,02% |
| java.time.LocalDate.ofEpochDay(long) | 24 bytes | 1,25 GB | 5,51% |
| java.time.LocalTime.ofNanoOfDay(long) | 24 bytes | 1,02 GB | 4,46% |
| java.time.LocalTime.create(int, int, int, int) | 24 bytes | 1,02 GB | 4,46% |

JMC / Flight Recorder

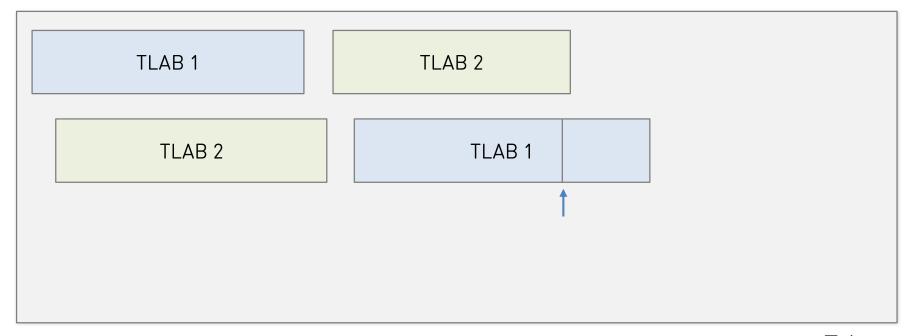


| Allocation Profile | | | |
|---|---------------------|-----------------|----------|
| Stack Trace | Average Object Size | Total TLAB size | Pressure |
| y perf.JavaTime.main(String[]) y perf.JavaTime.main(Stri | 37 bytes | 22,75 GB | 99,94% |
| perf.JavaTime.measureTimeObjects(int) | 37 bytes | 22,75 GB | 99,94% |
| java.time.LocalDateTime.now() | 37 bytes | 22,75 GB | 99,94% |
| > java.time.Clock.systemDefaultZone() | 39 bytes | 19,33 GB | 84,91% |
| java.time.LocalDateTime.now(Clock) | 24 bytes | 3,42 GB | 15,02% |
| java.time.LocalDateTime.ofEpochSecond(long, int, ZoneOffset) | 24 bytes | 3,42 GB | 15,02% |
| java.time.LocalDate.ofEpochDay(long) | 24 bytes | 1,25 GB | 5,51% |
| java.time.LocalTime.ofNanoOfDay(long) | 24 bytes | 1,02 GB | 4,46% |
| java.time.LocalTime.create(int, int, int, int) | 24 bytes | 1,02 GB | 4,46% |

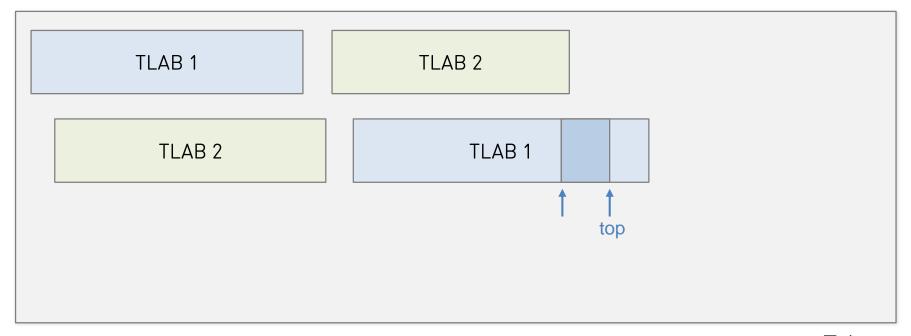


Overhead < 5%</pre>

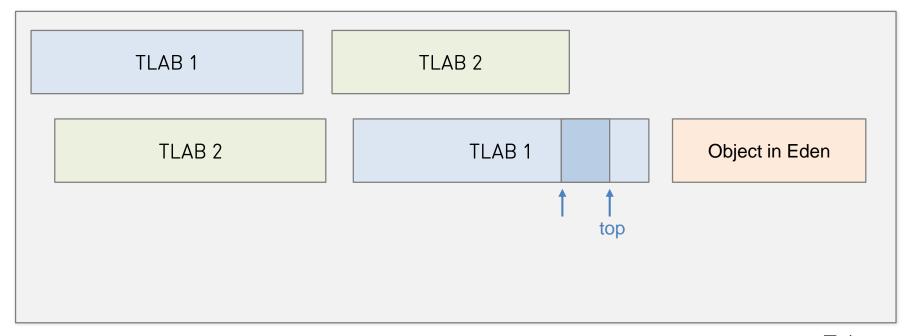




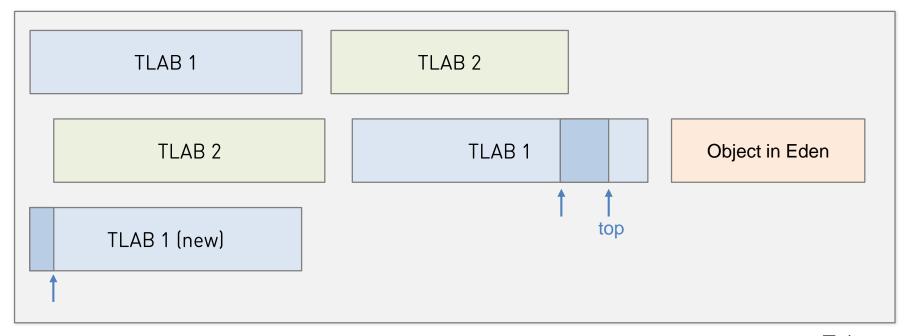












Allocation path

8

- Fast: inlined
 - Inside TLAB

- Slow: call to VM Runtime
 - Outside TLAB
 - Allocation of new TLAB



Allocation sampler in async-profiler

8

- 1. Intercept slow allocations in JVM runtime
- 2. Record allocated object + stack trace
- 3. PROFIT!

Allocation sampler in async-profiler

8

- 1. Intercept slow allocations in JVM runtime
- 2. Record allocated object + stack trace

3. PROFIT!

- Works with OpenJDK 7u40+
- Does not rely on commercial features

Appeared in JDK 11



JEP 331: Low-Overhead Heap Profiling



Demo

Contribute to async-profiler



- Try it out
- Submit issues
- Update documentation
- Send pull requests



Thank you

@AndreiPangin https://pangin.pro